

```

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import math
import seaborn as sns

import statsmodels.api as sm
from statsmodels.stats.outliers_influence import variance_inflation_factor
import statsmodels.formula.api as smf
from statsmodels.stats.anova import anova_lm

!pip install ISLP
from ISLP import load_data
from ISLP.models import (ModelSpec as MS ,summarize, poly)

```



 Show hidden output



- 8a. Use the `sm.OLS()` function to perform a simple linear regression with mpg as the response and horsepower as the predictor. Use the `summarize()` function to print the results.

```

Auto = load_data("Auto")
Auto.columns
X = X = pd.DataFrame({'intercept':np.ones(Auto.shape [0]) ,
'horsepower': Auto['horsepower']})
Y = Auto['mpg']
model = sm.OLS(Y, X)
results = model.fit()
summarize(results)

```

	coef	std err	t	P> t	
intercept	39.9359	0.717	55.660	0.0	
horsepower	-0.1578	0.006	-24.489	0.0	

- From the results summary we can see that there is a relationship between horsepower and mpg. There is a relationship between the two as the t critical values tells us that we can reject the null hypothesis of the predictor having a 0 magnitude effect on the outcome variable. However for every 10 unit increase in horsepower the mpg of the car is only predicted to decrease by 1.578 mpg which does show a weak affect that the predictor has on mpg. The relationship between the two is negative.

```

design = MS(['horsepower'])
design = design.fit(Auto)
X = design.transform(Auto)
model = sm.OLS(Y, X)
new_df = pd.DataFrame ( {'horsepower':[98]})
newX = design. transform (new_df)
newX
new_predictions = results. get_prediction (newX);
print(new_predictions.predicted_mean)
print(new_predictions.conf_int(alpha =0.05))

```

```
[24.46707715]
[[23.97307896 24.96107534]]
```

The predicted mpg at 98 horsepower is 24.467 mpg. With 95% confidence we can say the true value is between 23.973 and 24.961.


8b - Plot the response and the predictor in a new set of axes `ax`. Use the `ax.axline()`

- ✓ method or the `abline()` function defined in the lab to display the least squares regression line.

```
def abline(ax , b, m, *args , ** kwargs):
    "Add a line with slope m and intercept b to ax"
    xlim = ax.get_xlim()
    ylim = [m * xlim [0] + b, m * xlim [1] + b]
    ax.plot(xlim , ylim , *args , ** kwargs)
```


```
ax = Auto.plot.scatter('horsepower', 'mpg')
abline(ax ,
        results.params [0],
        results.params [1],
        'r--',
        linewidth =3)
```

```
<ipython-input-498-de519c427683>:11: FutureWarning: Series.__getitem__ treating keys as positions is deprecated. In a future
results.params [0],
<ipython-input-498-de519c427683>:12: FutureWarning: Series.__getitem__ treating keys as positions is deprecated. In a future
results.params [1],
```




- ✓ 8c - plot diagnostic plots

```
ax = plt.subplots(figsize =(8 ,8))[1]
ax.scatter(results.fittedvalues , results.resid)
ax.set_xlabel ('Fitted value ')
ax.set_ylabel ('Residual ')
ax.axhline (0, c='k', ls='--');
```



```
infl = results.get_influence()
ax = plt.subplots(figsize=(8,8))[1]
ax.scatter(np.arange(X.shape[0]), infl.hat_matrix_diag)
ax.set_xlabel('Index')
ax.set_ylabel('Leverage')
np.argmax(infl.hat_matrix_diag)
```


↶ 115



Residual seems to grow at a higher rate for increase levels of fitted values which can indicate inaccuracy in the relationship and our prediction model. At lower index levels there is also higher leverage which could also signify more inaccuracy or undefitting of our line.

✓ 9a - Produce a scatterplot matrix which includes all of the variables in the data set.


```
sns.pairplot(Auto)
```

 <seaborn.axisgrid.PairGrid at 0x79b17c76d110>





- 9b - Compute the matrix of correlations between the variables using the `DataFrame.corr()` method.

Auto.corr()



	mpg	cylinders	displacement	horsepower	weight	acceleration	year	origin
mpg	1.000000	-0.777618	-0.805127	-0.778427	-0.832244	0.423329	0.580541	0.565209
cylinders	-0.777618	1.000000	0.950823	0.842983	0.897527	-0.504683	-0.345647	-0.568932
displacement	-0.805127	0.950823	1.000000	0.897257	0.932994	-0.543800	-0.369855	-0.614535
horsepower	-0.778427	0.842983	0.897257	1.000000	0.864538	-0.689196	-0.416361	-0.455171
weight	-0.832244	0.897527	0.932994	0.864538	1.000000	-0.416839	-0.309120	-0.585005
acceleration	0.423329	-0.504683	-0.543800	-0.689196	-0.416839	1.000000	0.290316	0.212746
year	0.580541	-0.345647	-0.369855	-0.416361	-0.309120	0.290316	1.000000	0.181528
origin	0.565209	-0.568932	-0.614535	-0.455171	-0.585005	0.212746	0.181528	1.000000

- 9c - Use the `sm.OLS()` function to perform a multiple linear regression with mpg as the response and all other variables except name as the predictors.

- I tried doing this method the lab did in the book but have an issue as Pandas and this method seem incompatible even though the book says it is possible. So I am doing the following method I found online.


```
formula = 'mpg ~ cylinders + displacement + horsepower + weight + acceleration + year + origin'
multi_model = smf.ols(formula, data=Auto)
results_multi = multi_model.fit()
anova_lm(results_multi)
```

	df	sum_sq	mean_sq	F	PR(>F)
cylinders	1.0	14403.083079	14403.083079	1300.683788	2.319511e-125
displacement	1.0	1073.344025	1073.344025	96.929329	1.530906e-20
horsepower	1.0	403.408069	403.408069	36.430140	3.731128e-09
weight	1.0	975.724953	975.724953	88.113748	5.544461e-19
acceleration	1.0	0.966071	0.966071	0.087242	7.678728e-01
year	1.0	2419.120249	2419.120249	218.460900	1.875281e-39
origin	1.0	291.134494	291.134494	26.291171	4.665681e-07
Residual	384.0	4252.212530	11.073470	NaN	NaN

From the data we can see that most of the variables besides acceleration seem to have a strong correlation to the multi model. we fail to reject a 0 magnitude effect from acceleration in this model. Cylinders, Displacement, Weight, Horsepower and Year seemed to have the most statistical significance in impacting mpg. The coefficient means that for a car that is newer by a year (has an 1 increase in the year value) the mpg saw an increase by about 0.58 mpg.


9d - Produce some of diagnostic plots of the linear regression fit as described in the lab. Comment on any problems you see with the fit. Do the residual plots suggest any unusually large outliers? Does the leverage plot identify any observations with unusually high leverage?

```
ax = plt.subplots(figsize=(8,8))[1]
ax.scatter(results_multi.fittedvalues , results_multi.resid)
ax.set_xlabel('Fitted value ')
ax.set_ylabel('Residual ')
ax.axhline(0, c='k', ls='--');
```



```
infl = results_multi.get_influence()
ax = plt.subplots(figsize=(8,8))[1]
ax.scatter(np.arange(X.shape[0]), infl.hat_matrix_diag)
ax.set_xlabel('Index')
ax.set_ylabel('Leverage')
np.argmax(infl.hat_matrix_diag)
```


↶ 13



There was high residuals at lower and higher fitted values as well as higher leverage for the lower indexes which seemed like a big outlier. However the leverage seems rather consistent and not out of normal levels except for one outlier which had abnormally higher leverage than other values. This could influence the trend our model visualizes which could stray away from the true model.

9e - Fit some models with interactions as described in the lab. Do any interactions appear to be statistically significant?

```
formula_interaction = 'mpg ~ cylinders * displacement + horsepower * weight + acceleration * year + origin'
interaction_model = smf.ols(formula_interaction, data=Auto)
results_interaction = interaction_model.fit()
print(results_interaction.summary())
anova_lm(results_interaction, results_multi)
```



OLS Regression Results

=====						
Dep. Variable:	mpg	R-squared:	0.875			
Model:	OLS	Adj. R-squared:	0.871			
Method:	Least Squares	F-statistic:	266.1			
Date:	Thu, 06 Feb 2025	Prob (F-statistic):	4.42e-165			
Time:	04:14:39	Log-Likelihood:	-953.98			
No. Observations:	392	AIC:	1930.			
Df Residuals:	381	BIC:	1974.			
Df Model:	10					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]

Intercept	111.4537	18.485	6.030	0.000	75.109	147.798
cylinders	-0.0415	0.477	-0.087	0.931	-0.980	0.897
displacement	-0.0143	0.015	-0.939	0.348	-0.044	0.016
cylinders:displacement	0.0015	0.002	0.707	0.480	-0.003	0.006
horsepower	-0.2295	0.026	-8.848	0.000	-0.280	-0.178
weight	-0.0102	0.001	-11.380	0.000	-0.012	-0.008
horsepower:weight	5.151e-05	6.68e-06	7.714	0.000	3.84e-05	6.46e-05
acceleration	-6.9334	1.151	-6.023	0.000	-9.197	-4.670
year	-0.6434	0.240	-2.679	0.008	-1.116	-0.171
acceleration:year	0.0892	0.015	5.978	0.000	0.060	0.119
origin	0.6422	0.248	2.588	0.010	0.154	1.130
=====						
Omnibus:	40.489	Durbin-Watson:	1.603			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	86.168			
Skew:	0.563	Prob(JB):	1.94e-19			
Kurtosis:	5.002	Cond. No.	5.30e+07			
=====						

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
 [2] The condition number is large, 5.3e+07. This might indicate that there are strong multicollinearity or other numerical problems.

	df_resid	ssr	df_diff	ss_diff	F	Pr(>F)	
0	381.0	2982.899441	0.0	NaN	NaN	NaN	
1	384.0	4252.212530	-3.0	-1269.313089	38.208832	NaN	

Horsepower x Weight and Acceleration x Year were the only two that seemed to be statistically significant with cylinders x displacement having a low t value meaning we can't rule out the possibility of a 0 coefficient effect.

✓ 9f - Try a few different transformations of variables and comment findings.

```
Auto_transformed = Auto.copy()
```

```
Auto_transformed['log_horsepower'] = np.log(Auto_transformed['horsepower'])
Auto_transformed['sqrt_weight'] = np.sqrt(Auto_transformed['weight'])
Auto_transformed['displacement_sq'] = Auto_transformed['displacement'] ** 2
```

```
# Define formulas with transformed variables
```

```
transform_formulas = [
    'mpg ~ log_horsepower + weight + acceleration + year + origin',
    'mpg ~ horsepower + sqrt_weight + acceleration + year + origin',
    'mpg ~ cylinders + displacement_sq + horsepower + weight + acceleration + year + origin',
]
```

```
for formula in transform_formulas:
    model = smf.ols(formula, data=Auto_transformed)
    results = model.fit()
    print(f"Formula: {formula}")
    print(results.summary())
    print("\n" + "=" * 50 + "\n")
```



```

sqrt_weight -0.0892 0.053 -13.004 0.000 -0.793 -0.565
acceleration 0.0948 0.094 1.012 0.312 -0.089 0.279
year 0.7595 0.049 15.424 0.000 0.663 0.856
origin 0.9309 0.256 3.631 0.000 0.427 1.435
=====
Omnibus: 41.349 Durbin-Watson: 1.295
Prob(Omnibus): 0.000 Jarque-Bera (JB): 81.065
Skew: 0.602 Prob(JB): 2.49e-18
Kurtosis: 4.874 Cond. No. 3.91e+03
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
 [2] The condition number is large, 3.91e+03. This might indicate that there are strong multicollinearity or other numerical problems.

```

=====
Formula: mpg ~ cylinders + displacement_sq + horsepower + weight + acceleration + year + origin
OLS Regression Results

```

```

=====
Dep. Variable: mpg R-squared: 0.836
Model: OLS Adj. R-squared: 0.833
Method: Least Squares F-statistic: 279.9
Date: Thu, 06 Feb 2025 Prob (F-statistic): 1.57e-146
Time: 04:14:39 Log-Likelihood: -1006.7
No. Observations: 392 AIC: 2029.
Df Residuals: 384 BIC: 2061.
Df Model: 7
Covariance Type: nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	-13.7793	4.488	-3.070	0.002	-22.604	-4.954
cylinders	-0.7083	0.261	-2.710	0.007	-1.222	-0.194
displacement_sq	6.951e-05	1.07e-05	6.475	0.000	4.84e-05	9.06e-05
horsepower	-0.0425	0.014	-3.075	0.002	-0.070	-0.015
weight	-0.0064	0.001	-11.024	0.000	-0.008	-0.005
acceleration	0.0747	0.094	0.792	0.429	-0.111	0.260
year	0.7644	0.049	15.646	0.000	0.668	0.860
origin	1.3374	0.254	5.271	0.000	0.838	1.836

```

=====
Omnibus: 28.628 Durbin-Watson: 1.392
Prob(Omnibus): 0.000 Jarque-Bera (JB): 48.813
Skew: 0.473 Prob(JB): 2.51e-11
Kurtosis: 4.447 Cond. No. 1.93e+06
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
 [2] The condition number is large, 1.93e+06. This might indicate that there are strong multicollinearity or other numerical problems.

The transformed variables `log_horsepower` and `sqrt_weight` appear to provide a slightly improved fit without drastically changing the model's interpretation. The log transformation of horsepower shows that there may be a non linear relationship where an increase in horsepower has a diminishing effect on mpg while meaning smaller engines experience greater efficiency losses than larger ones. The square root transformation of weight showed a diminishing marginal impact of weight on fuel efficiency with the reinforcing the theory that heavier cars experience decreasing fuel efficiency losses as weight increases. These transformations help stabilize variance and may reduce multicollinearity.

- ✓ 13a - Using the normal() method of your random number generator create a vector, x, containing 100 observations drawn from a $N(0,1)$ distribution this represents a feature X.

```
x = np.random.normal(0,1,100)
print(x)
len(x)
```

```
[ -0.71896528 -0.34378447 -1.90126314 -0.7166678 -1.09313582 -0.53335378
  0.36020486 -1.19199502 -0.46485476  0.12300601 -0.41268306  2.07004512
 -1.85527234 -0.38975872 -0.14004807  0.31852385 -0.27715562 -0.74594951
  0.50877153 -0.12250746  1.56864964 -0.49433578  0.91988573  0.95886672
  1.44460813 -0.09978292 -0.00258054 -0.1969089 -0.14392027  0.94025683
  0.62119071  0.03390489 -0.04505755 -0.40027573 -0.80040501 -0.25120053
 -0.05990054  1.23542633 -0.73200465  1.62820874 -0.4633571  0.14388312
 -0.66630152 -0.10244666  0.35033646  1.39209807 -0.10236811 -0.105457
 -0.13125279  0.71040822  0.1587125  0.21055242 -0.26544538 -1.63809408
  0.69822819  1.12406972 -0.00695001  1.71591774 -0.72964932 -0.17485043
 -0.70993145 -0.7463968  0.26591789 -1.53825153 -1.52359683 -0.57863163
 -1.03810944  0.95843485  0.53822113  1.26119733  0.70161703  1.6376215
 -0.43924641  1.48685961 -0.50640188  1.27062862  1.23771299  0.74896209
  1.13169073  0.54482615  0.48997149 -0.87356548  0.26994964  0.56545755
  1.13469213  0.41607992  0.3431138  0.50452094 -0.14874546 -1.13711573
 -0.35623618  0.18688936  0.01947741  0.88116399 -0.0863606  1.07334577
 -0.166613 -1.4202214  0.81922492 -1.0800552 ]
100
```

- ✓ 13b - Using normal method create a vectore, eps, contining 100 obs drawn from a $N(0,0.25)$ dist.

```
eps = np.random.normal(0,0.5,100)
print(eps)
```

```
[ 0.16371882  0.14657842 -1.42959459  0.64502854 -0.38839181 -0.01002392
 -0.18145081  0.62021593  0.40304532 -0.37946853 -0.06142914  0.73286436
  0.27015559  0.03842355  0.18149251  0.26064104  0.49729375  0.31701974
 -0.20824801  0.94967055  1.10949837  0.41602695 -0.32726152  0.40577725
 -0.42074245 -0.33570135  0.45977608  0.18559324 -0.24101175 -0.56326801
 -0.14770939  0.38827517  0.2843148  0.15414563 -0.2368309 -0.68696089
  0.11772286  0.10077489  1.01509441 -0.1970753 -0.9119012  0.10638436
  0.4113634 -0.57063902 -1.15049578 -0.22045401 -0.33900044  0.45784901
 -0.10998022 -0.94086333 -0.70753865 -0.09717885 -0.28995288 -0.88726563
 -0.30499893  0.16120043  0.1872182 -0.18922199 -0.58117743 -0.93478573
  0.63918291 -0.7934551  0.11752694  0.26534985  0.02456851  0.07633705
  0.611251 -0.20234985  0.61669537  0.24871942  0.41560693 -0.48441315
 -0.54254096 -0.55736905 -0.79325342  0.53289726 -0.18547117  0.88306492
  0.89730986 -0.23323246 -0.62885952  0.88866358  0.66386174  0.41153587
 -0.19094446  0.48706131 -0.1288325  0.28864111 -0.37275201 -0.05064927
  0.78188331  0.54705889 -0.52219712 -0.42788504  0.35393427  0.81940282
 -1.08450814 -0.12125089  0.18746051 -0.03691782]
```

- ✓ 13c - using x and eps, generate a vector y according to model:

$$Y = -1 + 0.5X + \epsilon$$

what is the lenght of the vector? what are the values of b0 and b1 in this linear model

```
y = -1 + (x/2) + eps
print(y)
len(y)
```

```
[ -1.19576382 -1.02531381 -3.38022616 -0.71330536 -1.93495972 -1.27670081
 -1.00134838 -0.97578158 -0.82938206 -1.31796552 -1.26777067  0.76788692
 -1.65748058 -1.15645581 -0.88853153 -0.58009703 -0.64128407 -1.05595502
 -0.95386225 -0.11158318  0.89382319 -0.83114094 -0.86731866 -0.11478939
 -0.69843839 -1.38559281 -0.54151419 -0.91286121 -1.31297188 -1.09313959
 -0.83711404 -0.59477238 -0.73821398 -1.04599223 -1.63703341 -1.81256116
```

```

-0.9122274 -0.28151195 -0.35090792 -0.38297093 -2.14357975 -0.82167408
-0.92178736 -1.62186234 -1.97532755 -0.52440498 -1.3901845 -0.59487949
-1.17560662 -1.58565922 -1.6281824 -0.99190264 -1.42267557 -2.70631267
-0.95588484 -0.27676471 -0.81625681 -0.33126312 -1.94600209 -2.02221095
-0.71578282 -2.1666535 -0.74951411 -1.50377591 -1.7372299 -1.21297876
-0.90780372 -0.72313242 -0.11419407 -0.12068191 -0.23358455 -0.6656024
-1.76216417 -0.81393924 -2.04645435 0.16821157 -0.56661467 0.25754596
0.46315523 -0.96081939 -1.38387378 -0.54811916 -0.20116344 -0.30573535
-0.62359839 -0.30489873 -0.9572756 -0.45909842 -1.44712474 -1.61920713
-0.39623478 -0.35949643 -1.51245842 -0.98730305 -0.68924603 0.3560757
-2.16781464 -1.83136159 -0.40292703 -1.57694541]
100

```


The length of the vector is 100 meaning it has 100 values within it. B_0 is equal to -1 from the model we made above and b_1 is $1/2$

✓ 13d - Create a scatterplot displaying the relationship between x and y . Comment on what you observe.

```

plt.figure(figsize=(10, 10))
plt.scatter(x ,y )
plt.xlabel('x')
plt.ylabel('y')
plt.title('Scatterplot of x vs y')
plt.show()

```



This is a routine linear relationship with some noise around it. We can observe that the relationship looks almost perfectly linear and the data starts in the negative quadrant of the plot.

✓ 13e - Fit a least squares linear model to predict y using x. Comment on the model obtained. How do $\hat{\beta}_0$ and $\hat{\beta}_1$ compare to β_0 and β_1 ?

```
x = sm.add_constant(x)
model = sm.OLS(y,x)
results = model.fit()
print(results.summary())

b0_pred = results.params[0]
b1_pred = results.params[1]
print(f"Predicted b0: {b0_pred}")
print(f"Predicted b1: {b1_pred}")

print(f"Actual b0: -1")
print(f"Actual b1: 1/2")
```



OLS Regression Results

=====						
Dep. Variable:	y	R-squared:	0.440			
Model:	OLS	Adj. R-squared:	0.434			
Method:	Least Squares	F-statistic:	77.02			
Date:	Thu, 06 Feb 2025	Prob (F-statistic):	5.46e-14			
Time:	04:14:39	Log-Likelihood:	-77.747			
No. Observations:	100	AIC:	159.5			
Df Residuals:	98	BIC:	164.7			
Df Model:	1					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]

const	-0.9978	0.053	-18.700	0.000	-1.104	-0.892
x1	0.5457	0.062	8.776	0.000	0.422	0.669
=====						
Omnibus:	1.916	Durbin-Watson:	1.842			
Prob(Omnibus):	0.384	Jarque-Bera (JB):	1.671			
Skew:	-0.187	Prob(JB):	0.434			
Kurtosis:	2.488	Cond. No.	1.19			
=====						

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Predicted b0: -0.9977553842471643

Predicted b1: 0.5456980683513764

Actual b0: -1

Actual b1: 1/2

The actual and predicted coefficients were very close however not equal. there was an extremely small amount of error in both terms however they did capture the general trend of the data while be very clsoe to the true values.


13f - Display the least squares line on the scatterplot obtained in (d). Draw the

- ✓ population regression line on the plot, in a different color. Use the legend() method of the axes to create an **appropriate** legend.

```
pred_y = results.predict(x)
plt.figure(figsize=(10, 10))
plt.scatter(x[:,1], y)
plt.xlabel('x')
plt.ylabel('y')
plt.title('Scatterplot of y as function of x')

x_pop = np.linspace(x[:, 1].min(), x[:, 1].max(), 100)
y_pop = -1 + 0.5 * x_pop

plt.plot(x_pop, y_pop, color='blue', label='Population Regression Line')
plt.plot(x[:,1], pred_y, color='red', label='Least Squares Line')
plt.legend()
plt.show()
```



13g - Now fit a polynomial regression model that predicts y using x and x². Is there evidence that the quadratic term improves the model fit? Explain your answer.

```
x_squared = x[:, 1]**2
X_poly = np.column_stack((x, x_squared))
model_poly = sm.OLS(y, X_poly)
results_poly = model_poly.fit()
print(results_poly.summary())
```



OLS Regression Results

Dep. Variable:	y	R-squared:	0.441
Model:	OLS	Adj. R-squared:	0.429
Method:	Least Squares	F-statistic:	38.20
Date:	Thu, 06 Feb 2025	Prob (F-statistic):	5.82e-13
Time:	04:14:40	Log-Likelihood:	-77.699
No. Observations:	100	AIC:	161.4
Df Residuals:	97	BIC:	169.2
Df Model:	2		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	-0.9850	0.068	-14.453	0.000	-1.120	-0.850
x1	0.5484	0.063	8.689	0.000	0.423	0.674
x2	-0.0176	0.058	-0.304	0.762	-0.133	0.097

Omnibus:	1.990	Durbin-Watson:	1.835
Prob(Omnibus):	0.370	Jarque-Bera (JB):	1.660
Skew:	-0.168	Prob(JB):	0.436
Kurtosis:	2.466	Cond. No.	2.14

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

According to the summary statistics it seems that there is very little to assume that the x^2 variable contributed to our model with a t value that is very low not ruling out that it has a 0 magnitude effect on the y value.

13h - Repeat (a)–(f) after modifying the data generation process in such a way that there is less noise in the data. The model (3.39) should remain the same. You can do this by decreasing the variance of the normal distribution used to generate the error term ϵ in (b). Describe your results.

```
eps_reduced = np.random.normal(0,0.0625,100)
y = -1 + (x[:, 1]/2) + eps_reduced
model = sm.OLS(y,x)
results = model.fit()
print(results.summary())
```

```
x = sm.add_constant(x)
model = sm.OLS(y,x)
results = model.fit()
print(results.summary())
```

```
b0_pred = results.params[0]
b1_pred = results.params[1]
print(f"Predicted b0: {b0_pred}")
print(f"Predicted b1: {b1_pred}")
```

```
print(f"Actual b0: -1")
print(f"Actual b1: 1/2")
```



OLS Regression Results

Dep. Variable:	y	R-squared:	0.978
Model:	OLS	Adj. R-squared:	0.978
Method:	Least Squares	F-statistic:	4393.
Date:	Thu, 06 Feb 2025	Prob (F-statistic):	3.29e-83
Time:	04:14:40	Log-Likelihood:	131.94
No. Observations:	100	AIC:	-259.9
Df Residuals:	98	BIC:	-254.7
Df Model:	1		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
const	-0.9916	0.007	-151.289	0.000	-1.005	-0.979
x1	0.5062	0.008	66.276	0.000	0.491	0.521

Omnibus:	6.682	Durbin-Watson:	1.863
Prob(Omnibus):	0.035	Jarque-Bera (JB):	6.428
Skew:	0.474	Prob(JB):	0.0402
Kurtosis:	3.802	Cond. No.	1.19

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

OLS Regression Results

Dep. Variable:	y	R-squared:	0.978
Model:	OLS	Adj. R-squared:	0.978
Method:	Least Squares	F-statistic:	4393.
Date:	Thu, 06 Feb 2025	Prob (F-statistic):	3.29e-83

```

Time: 04:14:40 Log-Likelihood: 131.94
No. Observations: 100 AIC: -259.9
Df Residuals: 98 BIC: -254.7
Df Model: 1
Covariance Type: nonrobust

```

	coef	std err	t	P> t	[0.025	0.975]
const	-0.9916	0.007	-151.289	0.000	-1.005	-0.979
x1	0.5062	0.008	66.276	0.000	0.491	0.521
Omnibus:	6.682	Durbin-Watson:	1.863			
Prob(Omnibus):	0.035	Jarque-Bera (JB):	6.428			
Skew:	0.474	Prob(JB):	0.0402			
Kurtosis:	3.802	Cond. No.	1.19			

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Predicted b0: -0.9916142361101516

Predicted b1: 0.5062396638194264

Actual b0: -1

Actual b1: 1/2


```

pred_y = results.predict(x)
plt.figure(figsize=(10, 10))
plt.scatter(x[:,1], y)
plt.xlabel('x')
plt.ylabel('y')
plt.title('Scatterplot of y as function of x')

x_pop = np.linspace(x[:, 1].min(), x[:, 1].max(), 100)
y_pop = -1 + 0.5 * x_pop

plt.plot(x_pop, y_pop, color='blue', label='Population Regression Line')
plt.plot(x[:,1], pred_y, color='red', label='Least Squares Line')
plt.legend()
plt.show()

```



The reduced noise did make the ols line get closer to the true population line but that was due to lesser noise and data that strayed away from the true population line. This is the closest the OLS line has been to replacing the population line.

13i - Repeat (a)–(f) after modifying the data generation process in such a way that there is more noise in the data. The model (3.39) should remain the same. You can do this by increasing the variance of the normal distribution used to generate the error term ϵ in (b). Describe your results.

```
eps_increased = np.random.normal(0,0.75,100)
y = -1 + (x[:, 1]/2) + eps_increased
model = sm.OLS(y,x)
results = model.fit()
print(results.summary())
```

```
x = sm.add_constant(x)
model = sm.OLS(y,x)
results = model.fit()
print(results.summary())
```

```

b0_pred = results.params[0]
b1_pred = results.params[1]
print(f"Predicted b0: {b0_pred}")
print(f"Predicted b1: {b1_pred}")

print(f"Actual b0: -1")
print(f"Actual b1: 1/2")

```



OLS Regression Results

```

=====
Dep. Variable:          y      R-squared:          0.325
Model:                OLS      Adj. R-squared:       0.318
Method:             Least Squares  F-statistic:         47.13
Date:                Thu, 06 Feb 2025  Prob (F-statistic):    6.09e-10
Time:                04:14:40   Log-Likelihood:      -108.43
No. Observations:      100      AIC:                220.9
Df Residuals:          98      BIC:                226.1
Df Model:              1
Covariance Type:       nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	-1.0481	0.073	-14.453	0.000	-1.192	-0.904
x1	0.5802	0.085	6.865	0.000	0.412	0.748

```

=====
Omnibus:                0.363  Durbin-Watson:       2.151
Prob(Omnibus):          0.834  Jarque-Bera (JB):     0.528
Skew:                   -0.030  Prob(JB):             0.768
Kurtosis:               2.649  Cond. No.             1.19
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

OLS Regression Results

```

=====
Dep. Variable:          y      R-squared:          0.325
Model:                OLS      Adj. R-squared:       0.318
Method:             Least Squares  F-statistic:         47.13
Date:                Thu, 06 Feb 2025  Prob (F-statistic):    6.09e-10
Time:                04:14:40   Log-Likelihood:      -108.43
No. Observations:      100      AIC:                220.9
Df Residuals:          98      BIC:                226.1
Df Model:              1
Covariance Type:       nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	-1.0481	0.073	-14.453	0.000	-1.192	-0.904
x1	0.5802	0.085	6.865	0.000	0.412	0.748

```

=====
Omnibus:                0.363  Durbin-Watson:       2.151
Prob(Omnibus):          0.834  Jarque-Bera (JB):     0.528
Skew:                   -0.030  Prob(JB):             0.768
Kurtosis:               2.649  Cond. No.             1.19
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```

Predicted b0: -1.0481144669873066
Predicted b1: 0.5801988470511434
Actual b0: -1
Actual b1: 1/2

```


```

pred_y = results.predict(x)
plt.figure(figsize=(10, 10))
plt.scatter(x[:,1], y)
plt.xlabel('x')
plt.ylabel('y')
plt.title('Scatterplot of y as function of x')

x_pop = np.linspace(x[:, 1].min(), x[:, 1].max(), 100)
y_pop = -1 + 0.5 * x_pop

plt.plot(x_pop, y_pop, color='blue', label='Population Regression Line')
plt.plot(x[:,1], pred_y, color='red', label='Least Squares Line')
plt.legend()
plt.show()

```



The increased noise made the ols line less accurate to use to predict the population line. while the slopes were still close the intercept ended up moving lower. the increased noise has definitely made the ols underfit the model which will make residual errors between predictions and trye values larger.

- ✓ 13j - What are the confidence intervals for β_0 and β_1 based on the original data set, the noisier data set, and the less noisy data set? Comment on your result?

```
x = sm.add_constant(x)
model_original = sm.OLS(y, x)
results_original = model_original.fit()

conf_int_original = results_original.conf_int(alpha=0.05)

print(conf_int_original)
```

[[-1.19202563 -0.90420331]
[0.41248953 0.74790816]]

```

y_more_noise = -1 + (x[:, 1] / 2) + eps_increased
model_more_noise = sm.OLS(y_more_noise, x)
results_more_noise = model_more_noise.fit()

conf_int_more_noise = results_more_noise.conf_int(alpha=0.05)

print(conf_int_more_noise)

```

↗ `[[-1.19202563 -0.90420331]`
`[0.41248953 0.74790816]]`

```

y_less_noise = -1 + (x[:, 1] / 2) + eps_reduced
model_less_noise = sm.OLS(y_less_noise, x)
results_less_noise = model_less_noise.fit()

conf_int_less_noise = results_less_noise.conf_int(alpha=0.05)

print(conf_int_less_noise)

```

↗ `[[-1.00462128 -0.9786072]`
`[0.49108169 0.52139764]]`

From the confidence intervals we can see that the least varied interval was the least noisy while the one with the most variance was the data with more noise. both coefficients of b_0 and b_1 were closer to the true population values in the less noisy data and the least accurate or larger confidence intervals were from the noisy data.


15a - using the boston data set, For each predictor, fit a simple linear regression model to predict the response. Describe your results. In which of the models is there a statistically significant association between the predictor and the response?

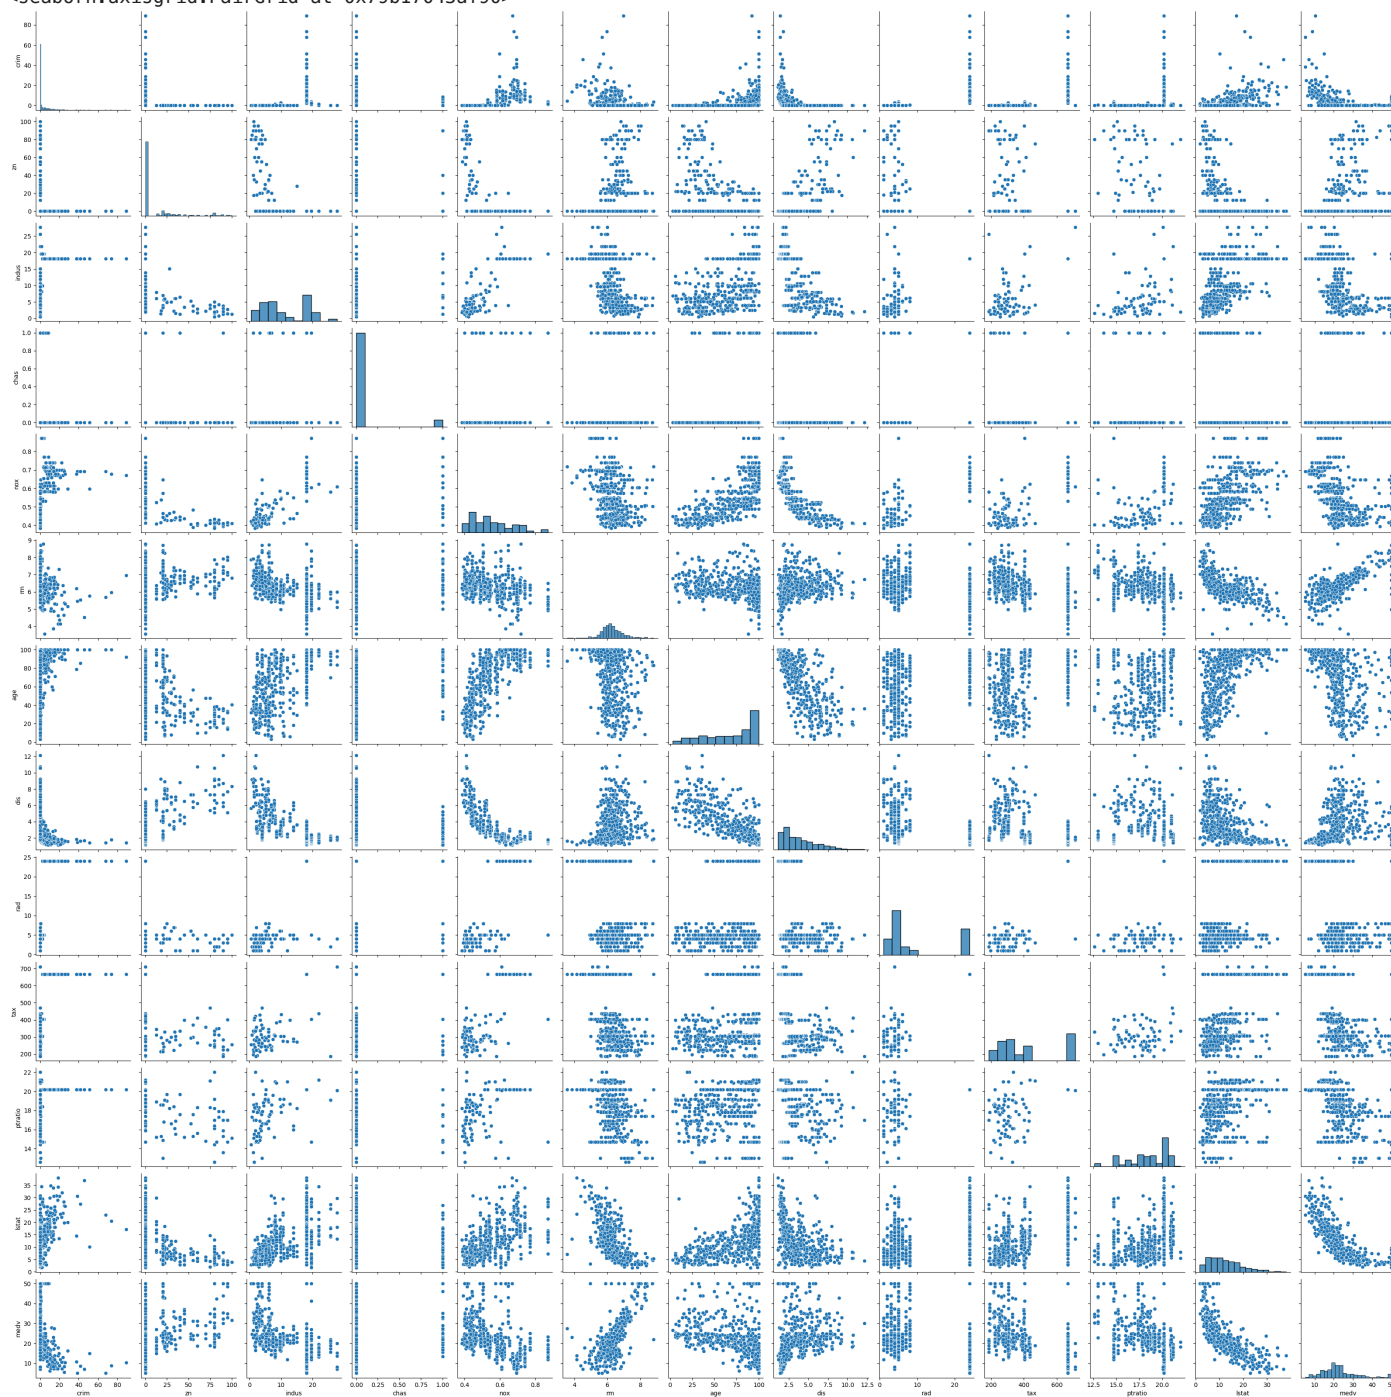
Create some plots to back up your assertions

```

from ISLP import load_data
boston_data = load_data('Boston')
boston = pd.DataFrame(boston_data)
sns.pairplot(boston)

```

 <seaborn.axisgrid.PairGrid at 0x79b17043af90>



```
predictors = boston.columns.drop('medv')
results = {}

for predictor in predictors:
    formula = f'medv ~ {predictor}'
    model = smf.ols(formula, data=boston).fit()
    results[predictor] = model

for predictor, model in results.items():
    print(f"Predictor: {predictor}")
    print(model.summary())
    print("-" * 50)
```




```

Dep. variable:      medv      R-squared:      0.544
Model:              OLS      Adj. R-squared:    0.543
Method:             Least Squares      F-statistic:    601.6
Date:               Thu, 06 Feb 2025      Prob (F-statistic): 5.08e-88
Time:              04:15:08      Log-Likelihood:   -1641.5
No. Observations:   506      AIC:              3287.
Df Residuals:       504      BIC:              3295.
Df Model:           1
Covariance Type:    nonrobust

```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	34.5538	0.563	61.415	0.000	33.448	35.659
lstat	-0.9500	0.039	-24.528	0.000	-1.026	-0.874
Omnibus:	137.043		Durbin-Watson:		0.892	
Prob(Omnibus):	0.000		Jarque-Bera (JB):		291.373	
Skew:	1.453		Prob(JB):		5.36e-64	
Kurtosis:	5.319		Cond. No.		29.7	

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.


```

for predictor in predictors:
    plt.figure(figsize=(8, 6))
    sns.regplot(x=predictor, y='medv', data=boston, ci=None, line_kws={'color': 'red'})
    plt.title(f'Simple Linear Regression: medv vs {predictor}')
    plt.xlabel(predictor)
    plt.ylabel('medv')
    plt.show()


```




Simple Linear Regression: medv vs crim




Simple Linear Regression: medv vs zn




Simple Linear Regression: medv vs indus






Simple Linear Regression: medv vs chas




Simple Linear Regression: medv vs nox




Simple Linear Regression: medv vs rm







Simple Linear Regression: medv vs age




Simple Linear Regression: medv vs dis




Simple Linear Regression: medv vs rad

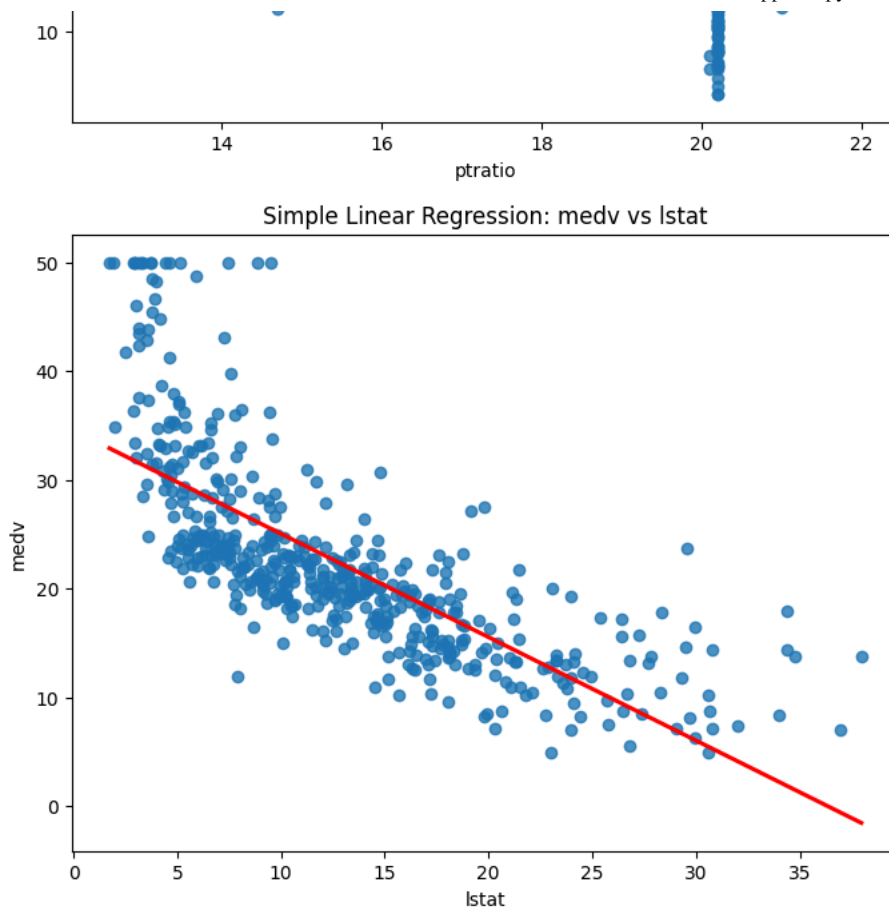


Simple Linear Regression: medv vs tax



Simple Linear Regression: medv vs ptratio





Some relationships that surfaced from the regression tables and the pairplots above came from predictors like Crm, Lstat, Age, Rm. These variables related to the crime rate within the suburb (crm), The percentage of people who are considered to be of lower socioeconomic status in the neighborhood (lstat), Age of the owner of the home (age), Avearge number of rooms in a dwelling. all these variable had some predictive relatiosbho whe setting median owner occupied house value as a dependent variable (medv).

15b - Fit a multiple regression model to predict the response using all of the

- ✓ predictors. Describe your results. For which predictors can we reject the null hypothesis $H_0 : \beta_j = 0$?

```
formula = 'medv ~ ' + ' + '.join(boston.columns.drop('medv'))
model = smf.ols(formula, data=boston).fit()
print(model.summary())
significant_predictors = [
    predictor for predictor in model.pvalues.index[1:]
    if model.pvalues[predictor] < 0.05 ]
print("\nSignificant Predictors (p-value < 0.05):")
print(significant_predictors)
```



OLS Regression Results

Dep. Variable:	medv	R-squared:	0.734
Model:	OLS	Adj. R-squared:	0.728
Method:	Least Squares	F-statistic:	113.5
Date:	Thu, 06 Feb 2025	Prob (F-statistic):	2.23e-133
Time:	04:15:11	Log-Likelihood:	-1504.9
No. Observations:	506	AIC:	3036.
Df Residuals:	493	BIC:	3091.
Df Model:	12		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
Intercept	41.6173	4.936	8.431	0.000	31.919	51.316
crim	-0.1214	0.033	-3.678	0.000	-0.186	-0.057
zn	0.0470	0.014	3.384	0.001	0.020	0.074
indus	0.0135	0.062	0.217	0.829	-0.109	0.136
chas	2.8400	0.870	3.264	0.001	1.131	4.549
nox	-18.7580	3.851	-4.870	0.000	-26.325	-11.191
rm	3.6581	0.420	8.705	0.000	2.832	4.484
age	0.0036	0.013	0.271	0.787	-0.023	0.030
dis	-1.4908	0.202	-7.394	0.000	-1.887	-1.095
rad	0.2894	0.067	4.325	0.000	0.158	0.421
tax	-0.0127	0.004	-3.337	0.001	-0.020	-0.005
ptratio	-0.9375	0.132	-7.091	0.000	-1.197	-0.678
lstat	-0.5520	0.051	-10.897	0.000	-0.652	-0.452

Omnibus:	171.096	Durbin-Watson:	1.077
Prob(Omnibus):	0.000	Jarque-Bera (JB):	709.937
Skew:	1.477	Prob(JB):	6.90e-155
Kurtosis:	7.995	Cond. No.	1.17e+04

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 1.17e+04. This might indicate that there are strong multicollinearity or other numerical problems.

Significant Predictors (p-value < 0.05):

['crim', 'zn', 'chas', 'nox', 'rm', 'dis', 'rad', 'tax', 'ptratio', 'lstat']

How do your results from (a) compare to your results from (b)? Create a plot displaying the univariate regression coefficients from (a) on the x-axis, and the multiple regression coefficients from (b) on the y-axis. That is, each predictor is displayed as a single point in the plot. Its coefficient in a simple linear regression model is shown on the x-axis, and its coefficient estimate in the multiple linear regression model is shown on the y-axis


```
univariate_coefs = {predictor: model.params[1] for predictor, model in results.items()}
multivariate_coefs = {predictor: model.params[predictor] for predictor in predictors}


x_coords = list(univariate_coefs.values())
y_coords = list(multivariate_coefs.values())
labels = list(univariate_coefs.keys())

plt.figure(figsize=(10, 8))
plt.scatter(x_coords, y_coords)

for i, label in enumerate(labels):
    plt.annotate(label, (x_coords[i], y_coords[i]), textcoords="offset points", xytext=(0, 10), ha='center')

plt.xlabel("Univariate Regression Coefficients")
plt.ylabel("Multiple Regression Coefficients")
plt.title("Comparison of Univariate and Multiple Regression Coefficients")
plt.grid(True)
plt.show()
```

 <ipython-input-526-0ef02e6e4b17>:1: FutureWarning: Series.__getitem__ treating keys as positions is deprecated. In a future



In part A i disregarded some valid coefficients due to my interpretation from the graphs and tables as single predictors looked increadibly inconsistent. but in the multi reg those variable did prove to be useful in addition to the other predictors.

15d- Is there evidence of non-linear association between any of the predictors and

✓ the response? To answer this question, for each predictor X, fit a model of the form

$$Y = \beta_0 + \beta_1 X + \beta_2 X^2 + \beta_3 X^3 + \epsilon.$$

```
for predictor in predictors:
    formula = f'medv ~ {predictor} + I({predictor}**2) + I({predictor}**3)'
    model = smf.ols(formula, data=boston).fit()
    print(f"Predictor: {predictor}")
    print(model.summary())
    print("-" * 50)
```

```
Model: OLS Adj. R-squared: 0.262
Method: Least Squares F-statistic: 60.91
Date: Thu, 06 Feb 2025 Prob (F-statistic): 1.35e-33
Time: 04:15:11 Log-Likelihood: -1761.7
No. Observations: 506 AIC: 3531.
Df Residuals: 502 BIC: 3548.
Df Model: 3
Covariance Type: nonrobust
```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	312.2864	152.487	2.048	0.041	12.695	611.878
ptratio	-48.6911	26.884	-1.811	0.071	-101.511	4.129
I(ptratio ** 2)	2.8400	1.564	1.816	0.070	-0.233	5.913
I(ptratio ** 3)	-0.0569	0.030	-1.892	0.059	-0.116	0.002

```
Omnibus: 98.389 Durbin-Watson: 0.735
Prob(Omnibus): 0.000 Jarque-Bera (JB): 200.314
Skew: 1.062 Prob(JB): 3.18e-44
Kurtosis: 5.235 Cond. No. 3.02e+06
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 3.02e+06. This might indicate that there are strong multicollinearity or other numerical problems.

Predictor: lstat

OLS Regression Results

```
Dep. Variable: medv R-squared: 0.658
Model: OLS Adj. R-squared: 0.656
Method: Least Squares F-statistic: 321.7
Date: Thu, 06 Feb 2025 Prob (F-statistic): 1.78e-116
Time: 04:15:11 Log-Likelihood: -1568.9
No. Observations: 506 AIC: 3146.
Df Residuals: 502 BIC: 3163.
Df Model: 3
Covariance Type: nonrobust
```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	48.6496	1.435	33.909	0.000	45.831	51.468
lstat	-3.8656	0.329	-11.757	0.000	-4.512	-3.220
I(lstat ** 2)	0.1487	0.021	6.983	0.000	0.107	0.191
I(lstat ** 3)	-0.0020	0.000	-5.013	0.000	-0.003	-0.001

```
Omnibus: 107.925 Durbin-Watson: 0.906
Prob(Omnibus): 0.000 Jarque-Bera (JB): 258.171
Skew: 1.088 Prob(JB): 8.69e-57
Kurtosis: 5.741 Cond. No. 5.20e+04
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.


[2] The condition number is large, 5.2e+04. This might indicate that there are strong multicollinearity or other numerical problems.

Crim, zn, indus, chas, rm, dis, rad, tax, lstat were the only predictors that had consistent low p-values on all terms which allowed for them to fit into non-linear regression. This intuitively makes sense for a lot of the variables where crime rates or even the number of rooms can have an accelerating or even diminishing returns for every increase which lowers the magnitude of the effect they have on medv. There is strong evidence both intuitive and in the models to prove the variables above have some non-linear relationship to the explanatory variable.

Part 3

✓ 1. Plot some relationships from the Homes data set and tell a story

```
homes_data = load_data('/homes2004')
homes = pd.DataFrame(homes_data)
sns.pairplot(homes)
```

 <seaborn.axisgrid.PairGrid at 0x79b169afc150>